

DM510 - Operating Systems

Project 3: Kernel Module

Jeff Gyldenbrand (jegyl16)
Supervisor: Daniel Merkle
Southern University of Denmark

April 23, 2018

Contents

1	Introduction	3
2	Design decisions	3
2.1	Devices	3
2.2	Buffers	3
3	Implementation	4
3.1	Devices	4
3.2	Buffers	4
3.3	Processes (sleep / awakening)	4
4	Test	5
5	Discussion	6
6	Conclusion	6
7	Appendix	7
7.1	Source-code	7

1 Introduction

This project is done in collaboration with Jonas Sørensen (joso216) and Simon D. Jørgensen (simjo16).

The goal of this project is to make a kernel module which implements a driver that exposes two character devices to user-space. This driver solves the producer-consumer problem, meaning the devices can both produce (write) and consume (read) data.

More specifically, two devices is to be implemented: dm510-0 and dm510-1. When a process writes to dm510-0, the data has to be stored in a buffer (buffer 1) that resides in kernel-space. In case this buffer is full, it must wait until another process reads from device dm510-1, and thereby emptying the buffer. When a process reads from dm510-0, the data is read from buffer 0, and now, if this buffer is empty, the process has to wait until another process writes to dm510-1. See fig:1 for graphical representation of this explanation.




fig:1

2 Design decisions

With the scull driver as inspiration source are large sections of the code, and features directly used from here, with minor changes. For more transparent code, the buffer is divided into its own header file: buffer.h, also, the ioctl-commands has its own header: dm510_ioctl-commands.h.

2.1 Devices

First the devices is setup and initialized, meaning, the device region is registered and memory is allocated for the buffers. Then a open-function handles whenever processes tries to open af device file, and a relase-function is invoked whenever a process closes a device-file. Furthermore a read-function handles the case where a process attempts to read from an already opened device. And lastly a write-function handles whenever a process writes to a device. Furthermore a syscall for ioctl is made to ..

2.2 Buffers

In the buffer header several functions is made. Firstly an initialization of a new buffer of given size with the buffer_init-function. A buffer_free_space-function to see how much

space we can write to. A `buffer_resize`-function to support scaling of the buffers. A `buffer_free`-function to free the buffer from memory. And A `buffer_write`-function and `buffer_read`-function to handle the read- and writes of the buffers.

3 Implementation

3.1 Devices

3.2 Buffers

The buffers consists of the struct `buffer`, and has a char pointer `rp` (read) and `wp` (write). The buffer is initialized in `buffer_init`-function with `kmalloc`, and setting both `rp` and `wp` to the beginning of the buffer.

To see how much space we can write to, this functions checks if read- and write-pointer is at same spot, hence, the whole buffer can be written to, else it calculates the write-space between them. See listing 1.

Listing 1: `buffer.h :: buffer_free_space`

```

1  if (head->rp == head->wp)
2      return head->size - 1;
3  return ((head->rp + head->size - head->wp) % head->size) - 1;

```

For the `buffer_resize`-function `memcpy` is invoked to copy complete memory blocks, and return a pointer to the new destination, and in this way, resize a buffer.

To free a buffer it is simply resized to zero. This is possible because if `krealloc`'s `new_size` is 0 and `p` is not a NULL pointer, the object pointed to is freed. See listing 2.

Listing 2: `buffer.h :: buffer_free`

```

1  kfree(head->buffer);
2  head->buffer = NULL;
3  return 0;

```

The `buffer_write`-function checks where the read- and write-pointer is relative to each other, to decide where to write a given input. This is the same principle for the `buffer_read`-function, to decide where to read from and to.

3.3 Processes (sleep / awakening)

When the module is initialized in the `dm510_init_module`-function a queue is initialized for each device. See listing 3.

Listing 3: `dm510_dev.c :: dm510_init_module`

```

1  for ( i = 0; i < DEVICE_COUNT; i++) {
2      init_waitqueue_head(&devices[i].inq);
3      init_waitqueue_head(&devices[i].outq);
4      mutex_init(&devices[i].mutex);
5      dprintf("Device(%d) = (%d, %d)", i, (i \% BUFFER_COUNT), ((i + 1)
        \% BUFFER_COUNT));

```

```

6         devices[i].read_buffer = buffers + (i \% BUFFER_COUNT);
7         devices[i].write_buffer = buffers + ((i + 1) \% BUFFER\_COUNT);
8         DEBUG_CODE(printk(""));
9         frame_device_setup(devices+i, global_device+i );
10    }

```

If the read- and write-pointer at any given time, in the dm510_read-function, is at the same spot, wait_event_interruptible is invoked. Adding the current thread to the queue, and calls the scheduler to schedule a new thread. See listing 4

Listing 4: dm510_dev.c :: dm510_read

```

1 if(wait_event_interruptible(dev->inq,(*rp != *wp))){
2     return rerror(-EAGAIN, "Reader was interrupted while sleeping.");
3 }

```

This is also handled in the buffer_write-function.

4 Test

Nine test has been made. The first test, as seen in the video at time-frame: 0:37 - 0:44, is a read- and write-test, with the program ./readwrite "I love kernel panic", to see that its possible to send a message and read it.

The second test is at time-frame: 0:59 - 1:08, to see when read from an empty buffer with blocking. Notice a ctrl-c is applied to interrupt.

The third test is at time-frame: 1:15 - 1:21, to see the same case as the second test, only with a non-blocking. This time an error is prompted: Error (-1).

The fourth test is at time-frame: 1:29 - 1:37, to see the case where more readers than allowed is applied. As shown, only 9 readers are allowed, the 10th reader prompts an -1 error: invalid pointer.

The fifth test is at time-frame: 1:45 - 1:48, to see the case where more than one writer is applied. Only one writer is legal, hence, the second writer prompts an error -1: invalid pointer.

The sixth test is at time-frame: 1:56 - 2:06, to see the case where a buffer that is full is written to, with blocking. Notice that ctrl-c is applied to interrupt.

The seventh test is at time-frame: 2:12 - 2:17, to see the same case as the sixth test, only with non-blocking. This time it prompts "Buffer is full"

The eighth test is at time-frame: 2:26 - 2:29, this test is the given moduletest. As shown it opens with success on r and w. Result 1 matches expected result 1, and result 2 matches expected result 2.

The final test is at time-frame: 2:37 - 3:12, to see that its possible to resize the buffer. In this case its resized to first "2024". Then to check that it is not possible to resize the buffer below used space, an string 'hello friend' is echoed into the device dm510-0, and its

attempted to resize the buffer to size "5". As shown an error is prompted: "Can't contract buffer size 5 below used space 14".

5 Discussion

Multiple processes to use a device simultaneously is not possible because only one writer per device can exist. If more writers is created, they will be rejected. Its possible to have more than one reader. But this is also handled.

6 Conclusion

A kernel module with a driver that exposes two character devises to user-space has been developed. It handles the producer-consumer problem in that, the devices can both write and read data, stored from buffers in kernel-space. Furthermore a process cant read from an empty buffer, but actively waits for another process to write to said buffer, before reading from it. And vice verca, a process cant write to a full buffer, so it wait to the buffer is emptied by another process. As shown in the test-section the major corner-cases is handled too. Every requirements has been met.

7 Appendix

7.1 Source-code

dm510_dev.c

```
1  /* Prototype module for second mandatory DM510 assignment */
2  #ifndef __KERNEL__
3  # define __KERNEL__
4  #endif
5  #ifndef MODULE
6  # define MODULE
7  #endif
8
9  //#define DEBUG
10 #include <linux/cdev.h>
11 #include <linux/module.h>
12 #include <linux/kernel.h>
13 #include <linux/fs.h>
14 #include <linux/uaccess.h>
15 /* Prototypes - this would normally go in a .h file */
16
17 //own header file for buffer, to make code readable
18 #include "buffer.h"
19 #include "dm510_ioctl_commands.h"
20 static int dm510_open( struct inode*, struct file* );
21 static int dm510_release( struct inode*, struct file* );
22 static ssize_t dm510_read( struct file*, char*, size_t, loff_t* );
23 static ssize_t dm510_write( struct file*, const char*, size_t, loff_t* );
24 long dm510_ioctl(struct file *filp, unsigned int cmd, unsigned long arg);
25
26 #define DEVICE_NAME "dm510_dev" /* Dev name as it appears in /proc/devices */
27 #define MAJOR_NUMBER 254
28 #define MIN_MINOR_NUMBER 0
29 #define MAX_MINOR_NUMBER 1
30
31 /* end of what really should have been in a .h file */
32
33 /* file operations struct */
34 static struct file_operations dm510_fops = {
35     .owner    = THIS_MODULE,
36     .read     = dm510_read,
37     .write    = dm510_write,
38     .open     = dm510_open,
39     .release  = dm510_release,
40     .unlocked_ioctl = dm510_ioctl
41 };
42
43 struct frame {
44     wait_queue_head_t inq, outq;    /* read and write queues */
45     struct buffer * read_buffer;
46     struct buffer * write_buffer;
47     int nreaders, nwriters;          /* number of openings for r/w */
48     struct fasync_struct *asynq_queue; /* asynchronous readers */
49     struct mutex mutex;              /* mutual exclusion semaphore */
50     struct cdev cdev;
51 };
```

```

52
53 static struct frame devices[DEVICE_COUNT];
54 static struct buffer buffers[BUFFER_COUNT];
55 static size_t max_processes = 10;
56
57 dev_t global_device = MKDEV(MAJOR_NUMBER,MIN_MINOR_NUMBER);
58
59 static int frame_device_setup(struct frame * dev, dev_t device){
60     cdev_init(&dev->cdev, &dm510_fops);
61     dev->cdev.owner = THIS_MODULE;
62     return cdev_add(&dev->cdev, device, 1);
63 };
64
65 #define BUFFER_DEFAULT_SIZE 4096
66
67 int dm510_init_module( void ) {
68     int i, result;
69     result = register_chrdev_region(global_device,DEVICE_COUNT,DEVICE_NAME);
70     if(result){
71         return error(result, "Failed to register chrdev_region.");
72     }
73     for (i = 0; i < BUFFER_COUNT; i++) {
74         result = buffer_init(buffers+i, BUFFER_DEFAULT_SIZE);
75         if(result < 0) return error(result, "Could not allocate memory
76             for buffer(%d).", i);
77     }
78     for ( i = 0; i < DEVICE_COUNT; i++) {
79         init_waitqueue_head(&devices[i].inq);
80         init_waitqueue_head(&devices[i].outq);
81         mutex_init(&devices[i].mutex);
82         dprintf("Device(%d) = (%d, %d)",i,(i % BUFFER_COUNT), ((i + 1) %
83             BUFFER_COUNT));
84         devices[i].read_buffer = buffers + (i % BUFFER_COUNT);
85         devices[i].write_buffer = buffers + ((i + 1) % BUFFER_COUNT);
86         DEBUG_CODE(printk(""));
87         frame_device_setup(devices+i, global_device+i );
88     }
89     return 0;
90 }
91
92 /* Called when module is unloaded */
93 void dm510_cleanup_module( void ) {
94     int i;
95     for(i = 0; i < DEVICE_COUNT ; i++){
96         if(devices[i].write_buffer) cdev_del(&devices[i].cdev);
97     }
98     for(i = 0; i < BUFFER_COUNT ; i++){
99         if(buffers[i].buffer) buffer_free(buffers+i);
100     }
101     unregister_chrdev_region(global_device,DEVICE_COUNT);
102 }
103
104
105 /* Called when a process tries to open the device file */

```



```

106 static int dm510_open( struct inode *inode, struct file *filp ) {
107     struct frame * dev;
108     dprintf("Open");
109     dev = container_of(inode->i_cdev, struct frame, cdev);
110     filp->private_data = dev;
111     if(mutex_lock_interruptible(&dev->mutex)){
112         return error(-ERESTARTSYS, "Mutex lock was interrupted.");
113     }
114
115     if (filp->f_mode & FMODE_READ){
116         if(dev->nreaders >= max_processes ){
117             mutex_unlock(&dev->mutex);
118             return error(-ERESTARTSYS, "Too many readers, only %d are
119                 allowed.", max_processes);
120         } else{
121             dev->nreaders++;
122         }
123     }
124     if (filp->f_mode & FMODE_WRITE){
125         if(dev->nwriters >= 1){
126             mutex_unlock(&dev->mutex);
127             return error(-ERESTARTSYS, "Amount of writers exceeded the
128                 allowed capacity of 1.");
129         }else{
130             dev->nwriters++;
131         }
132     }
133     mutex_unlock(&dev->mutex);
134     return nonseekable_open(inode, filp);
135 }
136
137
138 /* Called when a process closes the device file. */
139 static int dm510_release( struct inode *inode, struct file *filp ) {
140     struct frame * dev = filp->private_data;
141     dprintf("Release");
142     //scull_p_fasync(-1, filp, 0);
143     mutex_lock(&dev->mutex);
144     if (filp->f_mode & FMODE_READ && dev->nreaders)
145         dev->nreaders--;
146     if (filp->f_mode & FMODE_WRITE && dev->nwriters)
147         dev->nwriters--;
148     mutex_unlock(&dev->mutex);
149
150     return 0;
151 }
152
153
154 /* Called when a process, which already opened the dev file, attempts to read
155    from it. */
156 static ssize_t dm510_read( struct file *filp,
157     char *buf, /* The buffer to fill with data */
158     size_t count, /* The max number of bytes to read */
159     loff_t *f_pos ) /* The offset in the file */

```

```

159 {
160     dprintf("Read.");
161     struct frame * dev = filp->private_data;
162     char **rp = &dev->read_buffer->rp;
163     char **wp = &dev->read_buffer->wp;
164
165     if (mutex_lock_interruptible(&dev->mutex)){
166         return error(-ERESTARTSYS, "Mutex lock was interrupted by outside
167             source.");
168     }
169
170
171     while (*rp == *wp) {
172         mutex_unlock(&dev->mutex);
173         if (filp->f_flags & O_NONBLOCK){
174             return error(-EAGAIN, "File-pointer could not be blocked
175                 .");
176         }
177         dprintf("Read Sleeping.");
178         if(wait_event_interruptible(dev->inq,(*rp != *wp))){
179             return error(-EAGAIN, "Reader was interrupted while
180                 sleeping.");
181         }
182         dprintf("Reader Awoken.");
183         if(mutex_lock_interruptible(&dev->mutex)){
184             return error(-ERESTARTSYS, "Mutex lock was interrupted by
185                 outside source.");
186         }
187     }
188     count = buffer_read(dev->read_buffer,buf,count);
189     mutex_unlock (&dev->mutex);
190     wake_up_interruptible(&dev->outq);
191     dprintf("Read : Done.");
192     return count;
193 }
194
195 /* Called when a process writes to dev file */
196 static ssize_t dm510_write( struct file *filp,
197     const char *buf,/* The buffer to get data from */
198     size_t count, /* The max number of bytes to write */
199     loff_t *f_pos ) /* The offset in the file */
200 {
201
202     struct frame * dev = filp->private_data;
203     dprintf("Write.| wp = %d, rp = %d" ,
204         dev->write_buffer->wp - dev->write_buffer->buffer,
205         dev->write_buffer->rp - dev->write_buffer->buffer);
206     if (mutex_lock_interruptible(&dev->mutex))
207         return error(-ERESTARTSYS, "Mutex lock was interrupted by outside
208             source.");
209
210     if(count > buffers->size){
211         return error(-EMSGSIZE, "Message exceeded the size of the buffer
212             .");
213     }
214 }

```

```

209
210     while (buffer_free_space(dev->write_buffer) < count) {
211         //DEFINE_WAIT(wait);
212
213         mutex_unlock(&dev->mutex);
214         if (filp->f_flags & O_NONBLOCK){
215             return rerror(-EAGAIN, "File-pointer could not be blocked
216                 .");
217         }
218         if(wait_event_interruptible(dev->outq,
219             (buffer_free_space(dev->write_buffer) >= count))){
220             return rerror(-EAGAIN, "Writer was interrupted while
221                 sleeping.");
222         }
223         if (mutex_lock_interruptible(&dev->mutex))
224             return rerror(-ERESTARTSYS, "Mutex lock was interrupted by
225                 outside source.");
226     }
227     count = buffer_write(dev->write_buffer, (char*)buf, count);
228     dprintf("Write : Waking Reader. | wp = %d, rp = %d" ,
229         dev->write_buffer->wp - dev->write_buffer->buffer,
230         dev->write_buffer->rp - dev->write_buffer->buffer);
231     int i;
232     for(i = 0 ; i < DEVICE_COUNT ; i++){
233         wake_up_interruptible(&devices[i].inq);
234     }
235     //wake_up_interruptible(&dev->inq);
236
237     mutex_unlock (&dev->mutex);
238     dprintf("Write : Done.");
239     return count;
240 }
241
242 #define GET_BUFFER_SIZE 0
243 #define SET_BUFFER_SIZE 1
244 #define GET_MAX_NR_PROC 2
245 #define SET_MAX_NR_PROC 3
246
247 /* called by system call icotl */
248 long dm510_ioctl(
249     struct file *filp,
250     unsigned int cmd, /* command passed from the user */
251     unsigned long arg ) /* argument of the command */
252 {
253     switch(cmd){
254         case GET_BUFFER_SIZE:
255             return buffers->size;
256
257         case SET_BUFFER_SIZE:{
258             int i;
259             for(i = 0 ; i < BUFFER_COUNT ; i++){
260                 int used_space = buffers[i].size - buffer_free_space
261                     (buffers+i);
262                 if(used_space > arg) {

```

```

261         return rerror(-EINVAL, "Buffer(%d) has %lu
                amount of used space, cannot be reduced
                to size %lu.", i, used_space, arg);
262     }
263 }
264 for(i = 0 ; i < BUFFER_COUNT ; i++) {
265     buffer_resize(buffers+i,arg);
266 }
267 }
268 break;
269
270 case GET_MAX_NR_PROC:
271     return max_processes;
272
273 case SET_MAX_NR_PROC:
274     max_processes = arg;
275     break;
276
277 case GET_BUFFER_FREE_SPACE:
278     return buffer_free_space(buffers+arg);
279
280 case GET_BUFFER_USED_SPACE:
281     return buffers[arg].size - buffer_free_space(buffers+arg);
282 }
283
284
285     return 0; //has to be changed
286 }
287
288 module_init( dm510_init_module );
289 module_exit( dm510_cleanup_module );
290
291 MODULE_AUTHOR( "Jonas Ingerslev Soerensen, Jeff Gyldenbrand, Simon Dradrach
        Joergensen" );
292 MODULE_LICENSE( "GPL" );

```

dm510_ioctl_commands.h

```

1  #ifndef DM510_IOCTL_COMMANDS
2  #define DM510_IOCTL_COMMANDS
3
4  #define GET_BUFFER_SIZE 0
5  #define SET_BUFFER_SIZE 1
6  #define GET_MAX_NR_PROC 2
7  #define SET_MAX_NR_PROC 3
8  #define GET_BUFFER_FREE_SPACE 4
9  #define GET_BUFFER_USED_SPACE 5
10
11 #define DEVICE_COUNT 2
12 #define BUFFER_COUNT 2
13
14 #endif /* end of include guard: DM510_IOCTL_COMMANDS */

```

buffer.h

```

1  #ifndef DM510_BUFFER_H
2  #define DM510_BUFFER_H

```

```

3
4 #include <linux/slab.h>
5 #include <linux/errno.h>
6 #include "debug.h"
7
8 struct buffer{
9     char *buffer;
10    size_t size;
11    char *rp, *wp;
12    struct mutex mutex;
13 };
14
15 // return how much space we can write to. If read- and write-pointer is at same
16 // spot -> it means
17 // the whole buffer can be written to. Else, calculate the write-space between
18 // them.
19 size_t buffer_free_space(struct buffer * head){
20     if (head->rp == head->wp)
21         return head->size - 1;
22     return ((head->rp + head->size - head->wp) % head->size) - 1;
23 }
24
25 //In accordance with IOCTL declaraion section 1, the buffer shall support
26 // scaling.
27 int buffer_resize(struct buffer * head, size_t size){
28     void * pointer;
29     mutex_lock(&head->mutex);
30
31     pointer = kmalloc(size*sizeof(*head->buffer),GFP_KERNEL);
32
33     if(!pointer) return rerror(-ENOMEM);
34     if(head-> wp == head-> rp){
35         head->wp = head->rp = pointer;
36     } else if(head-> wp > head-> rp){
37         size = head->wp - head->rp;
38         memcpy(pointer,head->rp,size);
39         head->wp = pointer + size;
40         head->rp = pointer;
41     } else {
42         size = (head->buffer + head->size) - head->rp;
43         memcpy(pointer,head->rp,size);
44         head->rp = pointer;
45         memcpy(head->rp + size, head->buffer, head->wp - head->buffer);
46         head->wp = (head->rp + size) + (head->wp - head->buffer);
47     }
48     head->buffer = pointer;
49     kfree(head->buffer);
50     mutex_unlock (&head->mutex);
51     return 0;
52 }
53
54 // initialise a new buffer of given size.
55 int buffer_init(struct buffer * head, size_t size){

```

```

56     void * pointer;
57     DEBUG_CODE(if(head->buffer) dprintf("To prevent memory leaks, the buffer
        pointer should be NULL.")););
58
59     pointer = kmalloc(size * sizeof(*head->buffer),GFP_KERNEL);
60
61     if(!pointer) return rerror(-ENOMEM);
62
63     head->wp = head->rp = head->buffer = pointer;
64     head->size = size;
65
66     return 0;
67 }
68
69 // allocate memory for buffer
70 struct buffer * buffer(size_t size){
71     struct buffer * head = kmalloc(sizeof(*head), GFP_KERNEL);
72     buffer_init(head,size);
73     return head;
74 }
75
76 // free buffer by resizing to zero. This is possible because if
77 // krealloc's new_size is 0 and p is not a NULL pointer, the object pointed to is
    freed.
78 int buffer_free(struct buffer * head){
79     kfree(head->buffer);
80     head->buffer = NULL;
81     return 0;
82 }
83
84 size_t buffer_write(struct buffer * buf, char * seq, size_t size){
85     size_t new_size;
86     mutex_lock(&buf->mutex);
87     //dprintf("(%lu).wp : %lu" , (size_t)buf, (size_t)(buf->wp - buf->buffer))
        ;
88     if(buf->wp < buf->rp){
89         new_size = min((size_t)(buf->wp - buf->rp) - 1, size);
90         copy_from_user(buf->wp,seq,new_size);
91         buf->wp += new_size;
92
93     }else{
94         const size_t a = (buf->buffer + buf->size) - buf->wp;
95         const size_t b = (buf->rp - buf->buffer) % buf->size;
96         new_size = min(a,size);
97         copy_from_user(buf->wp,seq,new_size);
98         size -= new_size;
99         if(0 < size){
100             new_size = min(b,size);
101             buf->wp = buf->buffer;
102             copy_from_user(buf->wp,seq,new_size - 1);
103         }
104         buf->wp += new_size;
105     }
106     //dprintf("-> %lu\n" , (size_t)(buf->wp - buf->buffer));
107     mutex_unlock (&buf->mutex);
108     return new_size;

```

```

109 }
110 size_t buffer_read(struct buffer * buf, char * seq, size_t size){
111     size_t new_size = 0;
112     mutex_lock(&buf->mutex);
113     //dprintf("(%lu).rp : %lu" , (size_t)buf, (size_t)(buf->rp - buf->buffer))
114     ;
115     if(buf->rp < buf->wp){
116         new_size = min((size_t)(buf->wp - buf->rp), size);
117         copy_to_user(seq,buf->rp,new_size);
118         buf->rp += new_size;
119     } else {
120         const size_t a = (buf->buffer + buf->size) - buf->rp;
121         const size_t b = buf->rp - buf->buffer;
122
123         new_size = min(a,size);
124         copy_to_user(seq,buf->rp,new_size);
125         size -= new_size;
126         if(0 < size){
127             new_size = min(b,new_size);
128             buf->rp = buf->buffer;
129             copy_to_user(seq,buf->rp,new_size);
130         }
131         buf->rp += new_size;
132     }
133     //dprintf("-> %lu\n" , buf->rp - buf->buffer);
134     mutex_unlock (&buf->mutex);
135     return new_size;
136 }
137 }
138 #endif /* end of include guard: DM510_BUFFER_H */

```

debug.h

```

1  #ifndef DM510_DEBUG_H
2  #define DM510_DEBUG_H
3  #include <linux/slab.h>
4  #ifdef DEBUG
5  # define DEBUG_CODE(code) do {code} while (0);
6  #else
7  # define DEBUG_CODE(code) do {} while (0)
8  #endif
9  # define dprintf(...) DEBUG_CODE(printk("%s/%d : ",__FILE__,__LINE__); printk(
    __VA_ARGS__ ));)
10 # define rerror(errno,...) ({dprintf(KERN_ERR "Error (%d) : ",errno) ;
    DEBUG_CODE(printk(KERN_CONT __VA_ARGS__ ));); errno;})
11 #endif /* end of include guard: DM510_DEBUG_H */

```