

# DM510 - Operating Systems

## Project 4: File System

---

Jeff Gyldenbrand (jegyl16)  
Supervisor: Daniel Merkle  
Southern University of Denmark

---

April 23, 2018

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Design decisions</b>	<b>3</b>
<b>3</b>	<b>Implementation</b>	<b>4</b>
<b>4</b>	<b>Test</b>	<b>8</b>
<b>5</b>	<b>Discussion</b>	<b>8</b>
<b>6</b>	<b>Conclusion</b>	<b>8</b>
<b>7</b>	<b>Appendix</b>	<b>9</b>
7.1	Source-code . . . . .	9
7.1.1	header-files . . . . .	9
7.1.2	source-files . . . . .	12

# 1 Introduction

This project is done in collaboration with Jonas Sørensen (joso216) and Simon D. Jørgensen (simjo16).

The goal of this project is to implement a file-system in linux using FUSE (file system in userspace, which is an interface for linux that let users create their own filesystem without touching kernel space. More specific, the implementation should include directories and files, and a way to create, delete and list them. Furthermore a way to open, close, read and write files. Lastly the implementation should handle a way to show size, access- and modification time-stamps, to files.

# 2 Design decisions

The file system is implemented with inodes and implemented such, that they have the behavior of a b-plus three. Where the inodes is the buckets of the three. A struct-type *inode\_t* is developed, and has the elements *level*, *fill*, *keys* and *values*. The keys of the inode is highlighted in green, and their associated values is highlighted in grey. The inode is implented with size  $n = 4$ . See figure 1. If a node is a branch-node, it has  $n$  keys,  $n + 1$  values. If a node is a leaf-node, is has  $n$  keys,  $n$  values. In figure 1, a leaf node is illustrated.

The filesystem is one big file denoted as filesystem.fs. Again, see figure 1, where filesystem.fs consist of "0hello1world". In this example, the zero is the first key in the root inode. This key is located in memory at adress *a*. The key 0 has the associated value *hello*, which is located in memory at adress *b*. The next key is 1 with associated value *world*, located at the memory addresses, respectively, *c* and *d*.




fig:1

In figure 2, an example of the case, where a key  $-1$  with value *goodbye* is inserted, is illustrated. Here the pointers in the inode struct will be sorted and updated in according to the rules of the b-plus tree, hence key  $-1$  and its asociated value will be placed in the beginning of the inode.

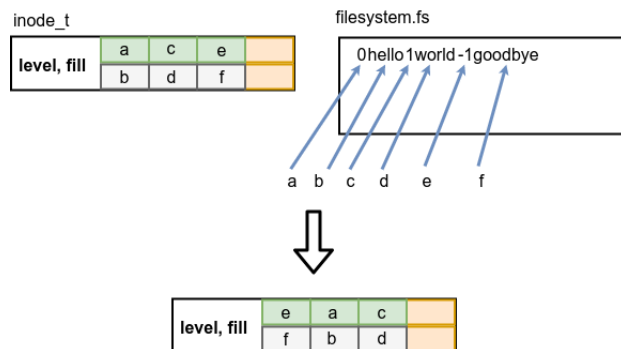



fig:2

To continue the example, a key 2 and its associated value *LOL* is inserted. However, the inode is full. Therefore the node is splitted into an  $\alpha$  and  $\beta$  inode, as seen in figure 3. A new root inode is created, with key  $c$  and associated pointers  $\alpha$  and  $\beta$ , which point to, respectively, the two leaf nodes  $\alpha$  and  $\beta$ . If a new key is to be inserted, and said key is less than  $c$ , said key would be assigned to the  $\alpha$  node, otherwise it would be assigned to the  $\beta$  node. Notice the root is now a branch-node, thus having  $n + 1$  values. This may seem odd, however, see it as this:  $c$  has a pointer to  $\alpha$  and  $\beta$ . If another key was inserted, this would have a pointer to  $\beta$  and  $\gamma$ , and yet another key would have pointers to  $\gamma$  and  $\delta$ .



Besides creating of new nodes, the implementation supports search for a given key, delete a key by replacing its pointer to zero, thus treating it as a NULL pointer, hence, now the spot is free. This is also how updates are handled; by deletion of a key, and insertion of a new one. All this will be elaborated in section 3.

### 3 Implementation

The examples in section 2 is just to illustrate the idea of the design. Here the keys are given as integers. The actual implementation of the b-plus three sorts lexographically, so the keys and values can be anything.

Several features have been developed to make the file system work. The datastructure for inode is implemented in inode.h and inode.c. The struct for inodes is shown in listing 1 below.

```


1 typedef struct{
2     file_ptr file_ptr;
3     size_t level;
4     size_t fill;
5     size_t keys[INODE_SIZE];
6     size_t values[1 + INODE_SIZE];
7     char data[];
8 } inode_t;

```

Listing 1: inode\_t

Due to the amplitude of functions created to succesfully make a filesystem, it would be overwhelming to cover them all. Therefore only the inode implementation will be covered in details here.

The function *node\_tree* is used to expand the b+ tree. This is done whenever a new root is created, and based on the new level, node tree will generate the same number of nodes. See figure 4, where the red nodes illustrates the newly expansion from root.



```

1 inode_t * node_tree(FILE *fp, size_t level){
2     inode_t n = {
3         .file_ptr.size = sizeof(n) - sizeof(file_ptr),
4         .level = level,
5         .fill = level > 0
6     };
7     if(level){
8         inode_t * result = node_tree(fp, level - 1);
9         n.values[0] = result->file_ptr.pos;
10        free(result);
11    }
12    return (inode_t *)append(fp, n.file_ptr.size, n.file_ptr.data);
13 }

```

Listing 2: node\_tree

The function *add* is responsively for inserting data. This is done with a key, provided to the b+ tree. See figure 5, case 1. This is an example where *add* is invoked with a key. Because the root has a node, *add\_recursive* is invoked. The number of *add\_recursive* invokes is based on the level, where e.g. level 2 would be equal to two recursive calls. In this case, the second recursive call is creating a new node due to its root has been splitted. This is returned to the first recursive call, and because its root has not been splitted, it returns null to the *add* function, which ultimaly return the new root.

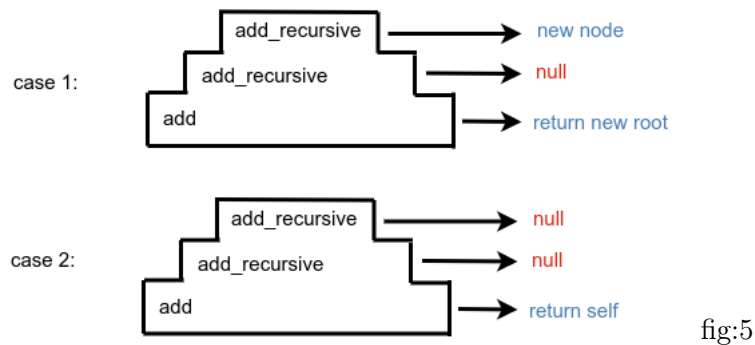
In case two both recursively calls return a null due to no roots has been splitted. Ultimately the *add* function will return itself, see listing 3. Here *add\_custom* is invoked with the key. Due to that *add\_custom\_key* is too large to display, please refer to appendix section 7.1.2 file *inode.c*, line 207 - 235, and 155 - 177 for *add\_recursive*.

```

1 inode_t * add(inode_t * root, char * key, size_t value_size, void * value){
2     return add_custom_key(root, strlen(key) + 1, key, value_size, value);
3 }

```

Listing 3: add



The *get* function firstly checks if a node is a branch-node, or leaf-node. If it is a branch-node it checks if the key fits between variables in the node, and moves it to child-node that fits the interval. If it is a leaf-node, it checks if the key equal to a variabel in the node.


```

1 ssize_t get(inode_t * node, char * key){
2     ssize_t index = find(node, key);
3     if(0 > index) {
4         return index;
5     }
6     ssize_t pos = node->values[index];
7     if(node->level && 0 <= pos){
8         node = struct_read(node->file_ptr.file, pos, *node);
9         pos = get(node, key);
10        free(node);
11    }
12    return pos;
13 }

```

Listing 4: get

The *tree\_key* function returns the leftmost key of a the b+ tree. See figure 6. In this example, nodes can have at most two elements. The root has elements 1 and 2. Then 3 and 4 is inserted, hence, the node is splitted into two nodes. A new root is created, and to assign a key to the new root, *tree\_key* is invoked. It searches untill it reaches the leftmost child, hence 3 is the new root key. Now 50, 100, 34 and 45 is inserted. Notice 34 is now the new key of yet another root.




```

1 file_ptr * tree_key(inode_t *node){
2     file_ptr * f;
3     if(node->level){
4         inode_t * n = struct_read(
5             node->file_ptr.file,
6             node->values[0],
7             *n
8         );
9         f = tree_key(n);
10        free(n);
11    } else{
12        f = get_string(
13            node->file_ptr.file,
14            node->keys[0]
15        );
16    }
17    return f;
18 }

```

Listing 5: tree\_key

The delete function resembles the get-function a lot in the way it finds the element to be deleted. In figure 7, element b is to be deleted. This is done by saving the memory address of b, then remove its pointer, hence now the spot where b was is free. Finally move elements c's pointer to the free spot.



```

1 ssize_t delete(inode_t * node, char * key){
2     ssize_t index = 0, pos = 0;
3     index = find(node, key);
4     pos = node->values[index];
5     if(node->level && 0 <= pos){
6         node = struct_read(node->file_ptr.file, pos, *node);
7         pos = delete(node, key);
8         free(node);
9     } else if(0 <= index){
10        memmove(
11            node->keys + index,
12            node->keys + index + 1,
13            (node->fill - index - 1) * sizeof(*node->keys)
14        );
15        node->keys[node->fill - 1] = 0;
16        memmove(
17            node->values + index,
18            node->values + index + 1,
19            (node->fill - index - 1) * sizeof(*node->values)
20        );
21        node->values[node->fill - 1] = 0;
22        node->fill--;
23        buf_write(&node->file_ptr);
24    } else{
25        pos = -1;
26    }
27    return pos;
28 }

```

Listing 6: delete

## 4 Test

Five test has been made to verify that the solution works correctly, and all requirements is meet.

**The first test**, as seen in the video at time-frame: 0:15 - 0:32, is to verify that its possible to create folders. First 'ls' is typed to show that the current working directory is empty. Then 'mkdir somefolder' is typed to create the folder, followed by another 'ls' command to verify it is created. To show that timestamps is also implemented, 'ls -alt' is typed too.

**Test two**, at time-frame 0:34 - 0:57, is to verify that its possible to create files, and add content to files. First 'cd somefolder' is typed to change directory into the newly created folder. Then 'echo "i like carrots" > rabbit.txt' is typed to create the file and its content. Then 'ls' to show the file. Then 'ls -alt' to verify the timestamp.

**Test three**, at time-frame 1:00 - 1:23, is to verify that its possible to read a file. With 'cat rabbit.txt' the file is read. To show its possible to append data 'echo ".. and beers" » rabbit.txt' is typed. Then 'cat rabbit.txt' is typed again, to read the updated file.

**Test four**, at time-frame 1:25 - 1:54, is to verify that its possible to delete files and folders. 'rm rabbit.txt' is typed to delete the file. Then 'ls' to verify its deleted. Then 'cd ..' to change to parent directory. Now 'rm -r somefolder/' is typed to delete the folder. Lastly 'ls' is typed to verify the deletion.

**Test five**, at time-frame 1:54 - 2:46, is to verify that its possible create multiple nested folders with files. A folder 'somefolder\_again' is created, inside this folder, a command to recursively create four folders with three subfolders in each, and one file in all folders, is typed. Then 'tree' command is typed to visualize the folders and files. Lastly all the folders are deleted, and 'tree' is typed again, to verify everything is deleted.

## 5 Discussion

Our implementation has a very high garanty to recover from a powerfailure due to the fact the file is updated frequently with the data. The only time when data would be unrecoverable is if the powerfailure would happen during a write to file, and even then, only the actual string currently being written would be lost, everything else would be stored in the file.

## 6 Conclusion

A file-system in linux using FUSE has been developed. Its possible to create and list both folders and files, and delete them again. Furthermore its possible to open-, read-, write- and close files. Lastly its possible to show access- and modification time-stamps of files. Every requirements has been met.



## 7 Appendix

### 7.1 Source-code

#### 7.1.1 header-files

file\_ptr.h

```
1 #ifndef DM510_FILE_PTR_H
2 #define DM510_FILE_PTR_H
3
4 #include <stdio.h>
5 #include <stdlib.h>
6 #include <string.h>
7
8 typedef struct {
9     FILE * file;
10    size_t pos;
11    size_t size;
12    char data[]; // <- Allows the expansion of the struct.
13 } file_ptr;
14
15 // The end of the filestream, returns the size of the filestream.
16 size_t fend(FILE *);
17
18 // Reads a buffer of a size from a position in a filestream.
19 file_ptr * buf_read(FILE *, size_t pos, size_t size);
20
21 // Writes a file pointer.
22 size_t buf_write(const file_ptr *);
23
24 // Appends a buffer at the end of the filestream.
25 file_ptr * append(FILE *, size_t size, const char * buf);
26
27 // Appends a file_ptr to the end of the filestream. OBS : updates the file
    pointers position.
28 int append_file_ptr(file_ptr * fp);
29
30 // Creates a new file_ptr, copies the buffer into itself.
31 file_ptr * new_file_ptr(FILE *fs, size_t pos, size_t size, const char * buf)
    ;
32
33 #define struct_read(file, pos, value) (typeof(value)*) buf_read(file, pos,
    sizeof(value) - sizeof(file_ptr))
34 #define primitiv_read(file, pos, value) ({ file_ptr * fs = buf_read(file, pos,
    sizeof(value)); typeof(value) a ; memcpy(&a, fs->data, sizeof(a)); free(fs)
    ; a ;})
35 #define struct_write(value) buf_write((file_ptr *) value)
36 #define struct_append(file, value) (typeof(value)*) append(file, sizeof(value)
    - sizeof(file_ptr), (const char *)&value + sizeof(file_ptr))
37
38 #endif /* end of include guard: DM510_FILE_PTR_H */
```

## fstruct.h

```
1 #ifndef DM510_FSTRUCT_H
2 #define DM510_FSTRUCT_H
3 #include <time.h>
4 #include "inode_files/inode.h"
5
6 typedef struct{
7     inode_t inode;
8     mode_t type;
9     char data[];
10 } group_t;
11
12
13 typedef struct{
14     group_t group;
15     size_t size;
16     time_t access_time; /* Time of last access */
17     time_t modification_time; /* Time of last modification */
18 } file_t;
19
20 typedef struct{
21     file_ptr file_ptr;
22     char page[1 << 12];
23 } page_t;
24
25 #endif /* end of include guard: DM510_FSTRUCT_H */
```

## inode.h

```
1 #ifndef DM510_INODE_H
2 #define DM510_INODE_H
3 #include "file_ptr.h"
4 #define KEY_EXISTS -1
5 #define KEY_NOT_EXISTS -2
6 #define FULL_INODE -3
7 #define INODE_SIZE 4
8
9 typedef struct{
10     file_ptr file_ptr;
11     size_t level;
12     size_t fill;
13     size_t keys[INODE_SIZE];
14     size_t values[1 + INODE_SIZE];
15     char data[];
16 } inode_t;
17
18 //Initialise the N and B printf keywords.
19 void init_printf_inode_extension();
20
21 //Get a string from a given file and position
22 file_ptr * get_string(FILE *fs, size_t pos);
23
24 //Create a b+ tree of a given depth, OBS: only the left most part of the
    nodes are filled.
25 inode_t * node_tree(FILE *fs, size_t depth);
26
27 //Get the leftmost key of a given b+ tree.
28 file_ptr * tree_key(inode_t *);
29
30 //Add some data, with a key, to the given b+ tree. OBS: Should the root splt
    , the new root will returned from the function.
```

```

31 inode_t * add(inode_t *, char *, size_t, void *);
32
33 inode_t * add_custom_key(inode_t *, size_t, void *, size_t, void *);
34
35 //Get the data of a given key.
36 ssize_t get(inode_t * , char *);
37
38 //Delete a key and its accompanying data.
39 ssize_t delete(inode_t * , char *);
40
41 #define struct_add(node, key, st) add(node, key, sizeof(st), &st)
42 #define struct_add_fptr(node, key, st) add(node, key, sizeof(st) - sizeof(
    file_ptr), (char *)&st + sizeof(file_ptr))
43
44
45 #endif /* end of include guard: DM510_INODE_H */

```

lfs.h

```

1 #ifndef DM510_LFS_H
2 #define DM510_LFS_H
3
4 #include <fuse.h>
5 #include <errno.h>
6 #include <string.h>
7 #include <stdio.h>
8 #include <stdlib.h>
9 #include <stdio.h>
10
11 int lfs_getattr( const char *, struct stat * );
12 int lfs_readdir( const char *, void *, fuse_fill_dir_t , off_t , struct
    fuse_file_info * );
13 int lfs_mknod( const char *, mode_t, dev_t);
14 int lfs_mkdir( const char *, mode_t);
15 int lfs_unlink( const char *);
16 int lfs_rmdir( const char *);
17
18 //int lfs_truncate( const char *, off_t, struct fuse_file_info *fi );
19 int lfs_open( const char *, struct fuse_file_info * );
20 int lfs_read( const char *, char *, size_t, off_t, struct fuse_file_info * )
    ;
21 int lfs_release( const char *path, struct fuse_file_info *fi);
22 int lfs_write( const char *, const char *, size_t, off_t, struct
    fuse_file_info *);
23 int lfs_utime( const char *, struct utimbuf *buf);
24
25 static struct fuse_operations lfs_oper = {
26     .getattr = lfs_getattr ,
27     .readdir = lfs_readdir ,
28     .mknod = lfs_mknod ,
29     .mkdir = lfs_mkdir ,
30     .unlink = lfs_unlink ,
31     .rmdir = lfs_rmdir ,
32     .truncate = NULL,
33     .open = lfs_open ,
34     .read = lfs_read ,
35     .release = lfs_release ,
36     .write = lfs_write ,
37     .rename = NULL,
38     .utime = lfs_utime
39 };
40

```

```

41
42
43 #endif /* end of include guard: DM510_LFS_H */

```

### 7.1.2 source-files

file\_ptr.c

```

1 #include "file_ptr.h"
2
3 file_ptr * new_file_ptr(FILE * fp, size_t pos, size_t size, const char *
    data){
4     file_ptr * f = malloc(sizeof(*f) + size);
5     f->file = fp;
6     f->pos = pos;
7     f->size = size;
8     if(data) memcpy(f->data, (char*)data, size);
9     return f;
10 }
11
12 size_t fend(FILE * fp){
13     fseek(fp,0,SEEK_END);
14     return ftell(fp);
15 }
16
17 file_ptr * buf_read(FILE * fp, size_t pos, size_t size){
18     fseek(fp,pos,SEEK_SET);
19     file_ptr * f = new_file_ptr(fp,pos,size,NULL);
20     int written = fread(f->data, size, 1, fp);
21     if(!written) {
22         free(f);
23         return NULL;
24     }
25     return f;
26 }
27
28 size_t buf_write(const file_ptr * fp){
29     fseek(fp->file,fp->pos,SEEK_SET);
30     return fwrite(fp->data,fp->size,1,fp->file);
31 }
32
33 file_ptr * append(FILE *fp, size_t size, const char * data){
34     file_ptr *f = new_file_ptr(fp,fend(fp),size,data);
35     int written = fwrite(f->data,f->size,1,f->file);
36     if(!written) {
37         free(f);
38         return NULL;
39     }
40     return f;
41 }
42
43 int append_file_ptr(file_ptr * f){
44     f->pos = fend(f->file);
45     return fwrite(f->data,f->size,1,f->file);
46 }

```

## inode.c

```

1 #include "inode.h"
2 #include <string.h>
3 #include <printf.h>
4
5
6
7 void print_key(FILE *stream, const inode_t * bpt, size_t i){
8     if(bpt->level){
9         fprintf(stream, "%lu", bpt->keys[i] );
10    } else {
11        fprintf(stream, "(%lu)", bpt->keys[i] );
12        if(bpt->keys[i]){
13            file_ptr * string = get_string(bpt->file_ptr.file, bpt->keys[i]) ;
14            fprintf(stream, "%s", string->data );
15            free(string);
16        }
17    }
18 }
19
20 int print_arginfo (const struct printf_info *info, size_t n, int *argtypes){
21     /* We always take exactly one argument and this is a pointer to the
22        structure.. */
23     if (n > 0)
24         argtypes[0] = PA_POINTER;
25     return 1;
26 }
27
28 int print_inode(FILE *stream, const struct printf_info *info, const void *
29     const *args){
30
31     const inode_t * bpt = *((const inode_t **) (*args));
32     int length = 0 ;
33     size_t i;
34     length += fprintf(stream, "%lu /%lu -> %lu : [", bpt->level, bpt->fill, bpt
35         ->file_ptr.pos );
36     for (i = 0; i < INODE_SIZE - 1; i++) {
37         print_key(stream, bpt, i);
38         length += fprintf(stream, ", ");
39     }
40     print_key(stream, bpt, i);
41     length += fprintf(stream, "], [%lu", bpt->values[0]);
42     for (i = 1; i < INODE_SIZE + 1; i++) {
43         length += fprintf(stream, ", %lu", bpt->values[i]);
44     }
45     length += fprintf(stream, "]");
46     return length;
47 }
48
49 int print_bpr_recursive(FILE *stream, const inode_t * bpt, size_t index){
50     const void *const arg = &bpt;
51     int sum = print_inode(stream, NULL, &arg);
52     if(!bpt->level) return sum;
53     for (size_t i = 0; i < bpt->fill; i++) {
54         inode_t * n = struct_read(bpt->file_ptr.file, bpt->values[i], *n);
55         if(bpt->values[i]){
56             sum += fprintf(stream, "\n%s", (int)index + 1, " ");
57             sum += print_bpr_recursive(stream, n, index + 1);
58         }
59         free(n);
60     }
61 }

```

```

58     }
59     return sum;
60 }
61
62 int print_bpr(FILE *stream, const struct printf_info *info, const void *
    const *args){
63     const inode_t * bpt = *((const inode_t **) (*args));
64     return print_bpr_recursive(stream, bpt, 0);
65 }
66
67 void init_printf_inode_extension(){
68     register_printf_function ('N', print_inode, print_arginfo);
69     register_printf_function ('B', print_bpr, print_arginfo);
70 }
71
72 file_ptr * get_string(FILE *fp, size_t pos){
73     size_t length = 1;
74     fseek(fp, pos, SEEK_SET);
75     while(fgetc(fp)) length++;
76     return buf_read(fp, pos, length);
77 }
78
79 file_ptr * tree_key(inode_t *node){
80     file_ptr * f;
81     if(node->level){
82         inode_t * n = struct_read(
83             node->file_ptr.file,
84             node->values[0],
85             *n
86         );
87         f = tree_key(n);
88         free(n);
89     } else{
90         f = get_string(
91             node->file_ptr.file,
92             node->keys[0]
93         );
94     }
95     return f;
96 }
97
98 ssize_t match(inode_t * node, char * key){
99     const size_t size = node->fill - (node->level > 0);
100     size_t i = 0;
101     for (; i < size; i++) {
102         file_ptr * string = get_string(node->file_ptr.file, node->keys[i]);
103         int cmp = strcmp(key, string->data);
104         free(string);
105         if(0 > cmp){
106             return i;
107         } else if(0 == cmp){
108             return node->level ? i + 1 : KEY_EXISTS;
109         }
110     }
111     return i;
112 }
113
114 inode_t * split_inode(inode_t * src, long index){
115     inode_t * dest = calloc(1, sizeof(*dest));
116     const size_t fsize = src->fill / 2, csize=(src->fill + 1) / 2;
117     size_t size = src->level ? 1 : 0;

```

```

118 dest->fill = fsize;
119 dest->file_ptr = src->file_ptr;
120 dest->level = src->level;
121 src->fill = csize;
122 size_t type = sizeof(*src->keys);
123 memcpy(dest->keys, src->keys + csize, (fsize - size) * type);
124 memset(src->keys + fsize, 0, fsize * type);
125
126 type = sizeof(*src->values);
127 memcpy(dest->values, src->values + csize, (fsize - size) * type);
128 memset(src->values + csize, 0, fsize * type);
129 return dest;
130 }
131
132 void node_make_room(inode_t * node, long index){
133     size_t size = (node->fill - index - (node->level > 0)) * sizeof(*node->
        keys);
134     memmove(node->keys + index + 1, node->keys + index, size);
135
136     size = (node->fill - index) * sizeof(*node->values);
137     memmove(node->values + index + 1, node->values + index, size);
138 }
139
140 inode_t * node_tree(FILE *fp, size_t level){
141     inode_t n = {
142         .file_ptr.size = sizeof(n) - sizeof(file_ptr),
143         .level = level,
144         .fill = level > 0
145     };
146     if(level){
147         inode_t * result = node_tree(fp, level - 1);
148         n.values[0] = result->file_ptr.pos;
149         free(result);
150     }
151     return (inode_t *)append(fp, n.file_ptr.size, n.file_ptr.data);
152 }
153
154
155 inode_t * add_recursive(inode_t * node, file_ptr * key, file_ptr * value){
156     inode_t * out_node = node, *new_node = NULL, *right_node = NULL;
157     ssize_t index = match(node, key->data);
158     int new_item = 0, split = 0;
159
160     if(0 > index) return new_node;
161     if(node->level){
162         if(!node->values[index]){ // Make new inode tree.
163             right_node = node_tree(node->file_ptr.file, node->level);
164         } else {
165             right_node = struct_read(
166                 node->file_ptr.file,
167                 node->values[index],
168                 *right_node
169             );
170         }
171         value = (file_ptr *)add_recursive(right_node, key, value);
172         free(right_node);
173         if(value){ // Handle split.
174             key = tree_key((inode_t *)value);
175             split = INODE_SIZE < node->fill;
176             new_item = 1;
177         }

```

```

178 } else {
179     split = INODE_SIZE == node->fill;
180     new_item = 1;
181 }
182 if(split){
183     new_node = split_inode(node, index);
184     append_file_ptr(&new_node->file_ptr); // Write the old node.
185     buf_write(&node->file_ptr); // Write the new node.
186     if(index > node->fill){
187         out_node = new_node;
188     }
189 }
190 if(!node->level || new_item){
191     index = match(out_node, key->data);
192     if(index < out_node->fill){
193         node_make_room(out_node, index);
194     }
195     out_node->values[index + (node->level > 0)] = value->pos;
196     out_node->fill++;
197     out_node->keys[index] = key->pos;
198 }
199 struct_write(out_node);
200 return new_node;
201 }
202
203 inode_t * add(inode_t * root, char * key, size_t value_size, void * value){
204     return add_custom_key(root, strlen(key) + 1, key, value_size, value);
205 }
206
207 inode_t * add_custom_key(inode_t * root, size_t key_size, void * key,
208     size_t value_size, void * value){
209     if(!(root && key && value)) return NULL;
210     size_t size = sizeof(inode_t) - sizeof(file_ptr);
211
212     file_ptr * key_ptr = append(root->file_ptr.file, key_size, key);
213     file_ptr * value_ptr = append(root->file_ptr.file, value_size, value);
214
215     inode_t * right_root = add_recursive(root, key_ptr, value_ptr);
216     if(right_root){ // Root got split.
217         inode_t * new_root = (inode_t *)new_file_ptr(root->file_ptr.file, 0, size,
218             NULL);
219         memset(new_root->file_ptr.data, 0, size);
220         key_ptr = tree_key(right_root);
221         new_root->keys[0] = key_ptr->pos;
222         new_root->file_ptr = root->file_ptr;
223         append_file_ptr(&root->file_ptr);
224
225         new_root->level = root->level + 1;
226         new_root->values[0] = root->file_ptr.pos;
227         new_root->values[1] = right_root->file_ptr.pos;
228         new_root->fill = 2;
229
230         buf_write((file_ptr *)new_root);
231         free(right_root);
232         root = new_root;
233     }
234     free(key_ptr);
235     free(value_ptr);
236     return root;
237 }

```



```

237 ssize_t find(inode_t * node, char * key){
238     size_t i = 0;
239     size_t size = node->fill - (node->level > 0);
240     for (; i < size; i++) {
241         file_ptr * string = get_string(node->file_ptr.file, node->keys[i]);
242         int cmp = strcmp(key, string->data);
243         free(string);
244         if((node->level && 0 > cmp) || (!node->level && !cmp)){
245             return i;
246         }
247     }
248     return node->level ? i : KEY_NOT_EXISTS;
249 }
250
251 ssize_t get(inode_t * node, char * key){
252     ssize_t index = find(node, key);
253     if(0 > index) {
254         return index;
255     }
256     ssize_t pos = node->values[index];
257     if(node->level && 0 <= pos){
258         node = struct_read(node->file_ptr.file, pos, *node);
259         pos = get(node, key);
260         free(node);
261     }
262     return pos;
263 }
264
265 ssize_t delete(inode_t * node, char * key){
266     ssize_t index = 0, pos = 0;
267     index = find(node, key);
268     pos = node->values[index];
269     if(node->level && 0 <= pos){
270         node = struct_read(node->file_ptr.file, pos, *node);
271         pos = delete(node, key);
272         free(node);
273     } else if(0 <= index){
274         memmove(
275             node->keys + index,
276             node->keys + index + 1,
277             (node->fill - index - 1) * sizeof(*node->keys)
278         );
279         node->keys[node->fill - 1] = 0;
280         memmove(
281             node->values + index,
282             node->values + index + 1,
283             (node->fill - index - 1) * sizeof(*node->values)
284         );
285         node->values[node->fill - 1] = 0;
286         node->fill--;
287         buf_write(&node->file_ptr);
288     } else{
289         pos = -1;
290     }
291     return pos;
292 }

```

## lfs.c

```

1 #include "lfs.h"
2 #include "fstruct.h"
3
4 #define min(x,y) x < y ? x : y
5 #define max(x,y) x > y ? x : y
6
7
8 FILE * file_system;
9
10 typedef struct{
11     char* string;
12     char* end;
13 } split_path_t;
14
15 split_path_t split_path(const char * path, char * (*f)(const char *, int)){
16     split_path_t p;
17     const char * end = f(path + 1, '/');
18     const char * pad = path + ('/' == *path);
19     if(!end){
20         size_t size = strlen(pad);
21         p.end = calloc(1, size + 1);
22         strncpy(p.end, pad, size);
23         p.string = NULL;
24     } else {
25         p.string = calloc(1, end - pad + 1);
26         strncpy(p.string, pad, end - pad);
27         const size_t size = strlen(end);
28         p.end = calloc(1, size + 1);
29         strncpy(p.end, end + 1, size);
30     }
31     return p;
32 }
33
34 group_t * walk(group_t * group, const char *path){
35     if(!path) return NULL;
36
37     split_path_t sp = split_path(path, strchr);
38     ssize_t pos = get(&group->inode, sp.string ? sp.string : sp.end);
39     if(0 > pos){
40         free(sp.string);
41         free(sp.end);
42         return NULL;
43     }
44     group_t * g = struct_read(group->inode.file_ptr.file, pos, *g);
45     if(sp.string){
46         group = walk(g, sp.end);
47         free(g);
48     } else {
49         group = g;
50     }
51     free(sp.string);
52     free(sp.end);
53     return group;
54 }
55
56
57 int lfs_getattr(const char *path, struct stat *stbuf) {
58     int res = 0;
59     group_t * root = struct_read(file_system, 0, *root);

```

```

60
61 memset(stbuf, 0, sizeof(struct stat));
62 if( strcmp( path, "/" ) == 0 ) {
63     stbuf->st_mode = S_IFDIR | 0755;
64     stbuf->st_nlink = 2;
65 } else {
66     group_t * group = walk(root, path);
67     if( group ){
68         stbuf->st_mode = group->type | 0777;
69         stbuf->st_nlink = 2;
70         if( S_IFREG == group->type) {
71             stbuf->st_nlink = 1;
72             file_ptr * fp = &group->inode.file_ptr;
73             file_t * file = struct_read(fp->file, fp->pos, *file);
74             stbuf->st_size = file->size;
75             stbuf->st_atime = file->access_time;
76             stbuf->st_mtime = file->modification_time;
77         }
78         free(group);
79     } else {
80         res = -ENOENT;
81     }
82 }
83 free(root);
84 return res;
85 }
86
87
88 void lfs_list_dir(inode_t * node, void *buf, fuse_fill_dir_t filler){
89     if(node->level){
90         for (size_t i = 0; i < node->fill; i++) {
91             inode_t * n = (inode_t*)buf_read(
92                 node->file_ptr.file,
93                 node->values[i],
94                 node->file_ptr.size
95             );
96             lfs_list_dir(n, buf, filler);
97             free(n);
98         }
99     } else {
100         for (size_t i = 0; i < node->fill; i++) {
101             file_ptr * string = get_string(node->file_ptr.file, node->keys[i]);
102             printf("string(%lu) = %s\n", node->keys[i], string->data );
103             filler(buf, string->data, NULL, 0);
104             free(string);
105         }
106     }
107 }
108
109 int lfs_readdir( const char *path, void *buf, fuse_fill_dir_t filler, off_t
    offset, struct fuse_file_info *fi ) {
110
111     printf("readdir: (path=%s)\n", path);
112     filler(buf, ".", NULL, 0);
113     filler(buf, "..", NULL, 0);
114
115     group_t * root = (group_t *)buf_read(file_system, 0, sizeof(group_t) -
        sizeof(file_ptr));
116     group_t * node = walk(root, path);
117     if(node){
118         lfs_list_dir(&node->inode, buf, filler);

```

```

119     free(node);
120 } else{
121     lfs_list_dir(&root->inode, buf, filler);
122 }
123 free(root);
124 return 0;
125 }
126
127 int lfs_mknod( const char * path, mode_t mode, dev_t dev){
128     group_t * root = struct_read(file_system, 0, *root);
129     split_path_t sp = split_path(path, strchr);
130     group_t * spot = walk(root, sp.string);
131     if(spot){
132         free(root);
133         root = spot;
134     }
135     file_t file = {
136         .group.type = S_IFREG,
137         .size = 0,
138         .access_time = time(NULL),
139         .modification_time = time(NULL)
140     };
141     struct_add_fptr(&root->inode, sp.end, file);
142     free(sp.string);
143     free(sp.end);
144     return 0;
145 }
146
147 int lfs_mkdir( const char * path, mode_t mode){
148     group_t * root = struct_read(file_system, 0, *root);
149     split_path_t sp = split_path(path, strchr);
150     group_t * spot = walk(root, sp.string);
151     if(spot){
152         free(root);
153         root = spot;
154     }
155     group_t group = {.type = S_IFDIR};
156
157     struct_add_fptr(&root->inode, sp.end, group);
158     free(root);
159     free(sp.string);
160     free(sp.end);
161     return 0;
162 }
163 int lfs_unlink(const char * path){
164     group_t * root = (group_t *)buf_read(file_system, 0, sizeof(group_t) -
165         sizeof(file_ptr));
166     split_path_t sp = split_path(path, strchr);
167     group_t * spot = walk(root, sp.string);
168     if(spot){
169         free(root);
170         root = spot;
171     }
172     delete(&root->inode, sp.end);
173     free(sp.string);
174     free(sp.end);
175     return 0;
176 }
177 int lfs_rmdir( const char * path){
178     group_t * root = (group_t *)buf_read(file_system, 0, sizeof(group_t) -

```

```

179     sizeof(file_ptr));
180     split_path_t sp = split_path(path, strchr);
181     group_t * spot = walk(root, sp.string);
182     if(spot){
183         free(root);
184         root = spot;
185     }
186     delete(&root->inode, sp.end);
187     free(sp.string);
188     free(sp.end);
189     return 0;
190 }
191 int lfs_open( const char *path, struct fuse_file_info *fi ) {
192     return 0;
193 }
194
195 char * hash(size_t key){
196     char *string = calloc(1,9);
197     sprintf(string, "%08x", (unsigned int)key);
198     return string;
199 }
200
201 page_t * merge_page(inode_t ** node, size_t key){
202     char * key_hash = hash(key);
203     ssize_t pos = get(*node, key_hash);
204     if(0 > pos){
205         static page_t p;
206         *node = struct_add(*node, key_hash, p);
207         pos = get(*node, key_hash);
208     }
209     free(key_hash);
210     return struct_read((*node)->file_ptr.file, pos, page_t);
211 }
212
213 int lfs_read( const char *path, char *buf, size_t size, off_t offset, struct
214     fuse_file_info *fi ) {
215     group_t * root = struct_read(file_system, 0, *root);
216     group_t * spot = walk(root, path);
217     free(root);
218     file_ptr *fp = &spot->inode.file_ptr;
219     file_t * file = struct_read( fp->file, fp->pos, *file );
220     free(spot);
221
222     const size_t start = offset, end = start + size;
223     while(offset < end) {
224         const off_t local_offset = offset % sizeof(page_t);
225         const size_t local_size = min(end - offset, sizeof(page_t) -
226             local_offset);
227
228         inode_t * node = &file->group.inode;
229         page_t * page = merge_page(&node, offset / sizeof(page_t) );
230         memcpy(offset - start + buf, page->page + local_offset, local_size);
231         offset += local_size;
232         free(page);
233     }
234     file->access_time = time(NULL);
235     struct_write(file);
236     return size;
237 }

```

```

237 int lfs_write( const char *path, const char *buf, size_t size, off_t offset,
238               struct fuse_file_info *fi){
239     group_t * root = struct_read(file_system, 0, *root);
240     group_t * spot = walk(root, path);
241     free(root);
242     file_ptr *fp = &spot->inode.file_ptr;
243     file_t * file = struct_read( fp->file, fp->pos, *file );
244     free(spot);
245
246     const size_t start = offset, end = start + size;
247     while(offset < end) {
248         const off_t local_offset = offset % sizeof(page_t);
249         const size_t local_size = min(end - offset, sizeof(page_t) -
250                                     local_offset);
251
252         inode_t * node = &file->group.inode;
253         page_t * page = merge_page(&node, offset / sizeof(page_t) );
254         memcpy(page->page + local_offset, offset - start + buf, local_size);
255
256         offset += local_size;
257         file->size += local_size;
258
259         struct_write(page);
260         free(page);
261     }
262     file->modification_time = time(NULL);
263     struct_write( file );
264     return size;
265 }
266
267 int lfs_release(const char *path, struct fuse_file_info *fi) {
268     return 0;
269 }
270
271 int lfs_utime( const char * path, struct utimbuf *buf){
272     return 0;
273 }
274
275 FILE * merge(char * file, char * mode){
276     FILE * fp = fopen(file, mode);
277     if(!fp) {
278         fp = fopen(file, "w");
279         fclose(fp);
280         fp = fopen(file, mode);
281     }
282     return fp;
283 }
284
285 int main( int argc, char *argv[] ) {
286     init_printf_inode_extension();
287     file_system = merge("filesystem.fs", "r+b");
288     size_t size = sizeof(group_t) - sizeof(file_ptr);
289     group_t * root = (group_t *)buf_read( file_system, 0, size);
290     if(!root){
291         printf("No root!\n");
292         group_t empty = {.type = S_IFDIR};
293
294         root = struct_append(file_system, empty);
295     }
296     printf("%B\n", root );
297     fuse_main( argc, argv, &lfs_oper );

```

```
296     return 0;  
297 }
```